

JAY RODRIGUEZ

Machine Learning Engineer

Hollywood, FL · jayalexander1127@gmail.com · 954-203-6334 · [linkedin.com/in/jay-rod](https://www.linkedin.com/in/jay-rod) · jayrodriguez.vercel.app · github.com/JayR1031

SUMMARY

ML engineer and MS Computer Science candidate (Northeastern, 2027) who builds and evaluates LLM systems for regulated domains, where a wrong answer has consequences. Shipped a zero-to-one enterprise AI platform with production RAG on Azure AI Foundry at diGenius AI, and founded Puente AI, whose deployed trade-compliance RAG system is the testbed for research on cross-lingual LLM compliance behavior (AlgoVerse AI Research, Fall 2026). Six years in enterprise sales to compliance-driven buyers before moving to engineering. Fluent in English and Spanish.

EXPERIENCE

AI Engineer | diGenius AI

Miami, FL (Remote) · May 2026 – Present

Restaurant-operations AI platform for multi-unit operators

- Architected a zero-to-one enterprise AI monorepo supporting 4 core applications (Portal, Buying, Brain, Cash Recon), with 500+ automated deployments to Azure Container Apps and Vercel through continuous integration.
- Designed and deployed secure retrieval-augmented generation (RAG) pipelines on LangChain and Azure AI Foundry over 3,000+ proprietary corporate records, across a Python (62%) and TypeScript (22%) codebase.
- Implemented CI/CD security gates with Trivy and Gitleaks on every pull request, blocking 32 vulnerable dependencies before production.

Founder & CEO | Puente AI

Miami, FL · May 2026 – Present

Trade-compliance and payment-routing platform for the US–Latin America corridor (Delaware C-Corp)

- Built and deployed a retrieval-augmented generation system on LanceDB and Google Cloud Run that answers customs and trade-compliance questions for brokers and shippers.
- Use the deployed system as a live research testbed for measuring how LLM compliance behavior degrades across English, Spanish, and Portuguese in regulated settings.

AI Evaluator (Contract) | Mercor

Remote · 2026 – Present

- Human evaluation of frontier-model outputs across generalist, healthcare, and power-user task domains: rubric-based scoring and preference judgments used for model training and evaluation.

Enterprise Sales & Account Management | Aegis IT · Uber Eats · Leverage AI · goTRG

Miami / NY · 2019 – 2026

- Sold managed IT and cybersecurity solutions to law firms, healthcare providers, and financial institutions, scoping deployments against HIPAA and SOC 2 requirements (Aegis IT Solutions, 2025–2026).
- Managed a \$3.5M portfolio of Southeast restaurant accounts, onboarding 60+ strategic locations and driving 18% growth in partner-funded promotions (Uber Eats, 2023–2025).
- Closed high-value deals averaging \$250K and launched a partnership worth \$3.5M in annual sales; 150% year-over-year growth on a \$2M book (goTRG, 2019–2022).

RESEARCH & PROJECTS

Cross-Lingual LLM Compliance Evaluation | AlgoVerse AI Research

Sep 2026 – Dec 2026 (in progress)

- Research question: how LLM compliance behavior fails across English → Spanish → Portuguese in regulated trade-compliance domains, positioned between recent results on over-compliance and cross-lingual representation drift. Target: workshop paper; continues as MS thesis (CS 7990).

PrivaSense MD — Privacy-First Medical Assistant | Team project, Northeastern

Nov 2025 – Dec 2025

github.com/CloudComputingMIA2025/team_awi

- Led frontend and cloud architecture on a 3-person team: Next.js UI, FastAPI backend, and VaultGemma (differentially private LLM) inference on ECS Fargate across 2 Availability Zones, with PHI masking before inference, encrypted RDS PostgreSQL, Cognito auth, and Terraform-managed infrastructure; load-tested to 100 concurrent users.

Multi-PDF RAG — Local, Privacy-Preserving Research Assistant | Independent

Nov 2025

github.com/JayR1031/multi-pdf-rag

- Engineered an end-to-end RAG pipeline (LangChain, ChromaDB) with fully local LLM inference and zero external API calls; ingestion, chunking, and embedding of multi-page PDFs with sub-2-second retrieval. Precursor to the Puente AI system.

Graphical Models for Intent Classification | CS 5100, Northeastern

2025

- Implemented Bayesian networks and Hidden Markov Models for customer-intent classification and text denoising.

EDUCATION

M.S. Computer Science (ML/AI concentration) | Northeastern University

Expected Aug 2027 · GPA 3.5

Coursework: Machine Learning, Artificial Intelligence, Algorithms, Cloud Computing, Data Structures, Object-Oriented Design. Graduate Student Advisory Board; HSF Scholar.

B.A.S. Supply Chain Management | Broward College

2021 · GPA 3.5

Process optimization, statistics, optimization & analytics, Six Sigma / Lean.

CERTIFICATIONS & LEADERSHIP

- Google Associate Cloud Engineer (2026).
- Miami Chapter Lead, The AI Collective (Jul 2025 – Present): run events and programs connecting founders, researchers, and operators across South Florida; built Airtable automations for chapter operations.
- Founding Chapter President, INIT @ Northeastern (Jul 2025 – Present): secured official university recognition in under 4 months and initial seed funding; launched 3 joint national-chapter workshops.

SKILLS

ML & NLP: Python, PyTorch, scikit-learn, NumPy, pandas, Hugging Face Transformers, LangChain, retrieval-augmented generation, LLM evaluation

Infrastructure: Docker, Terraform, Google Cloud (Cloud Run; certified), Azure AI Foundry & Container Apps, AWS (ECS Fargate, RDS), PostgreSQL, LanceDB, ChromaDB, FastAPI, Git, GitHub Actions, CI/CD security (Trivy, Gitleaks)

Mathematics: Linear algebra, probability, optimization

Languages: English (native), Spanish (fluent)